

```
import pandas as pd
import numpy as np

data=pd.read_csv("diabetes.csv")
```

data

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigr
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows x 9 columns

```
data.isnull().any()
```

```
Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction  False
Age               False
Outcome           False
dtype: bool
```

```
data.isna().sum()
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
data.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	I
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
data.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
x=data.iloc[:, :-1].values
y=data.iloc[:, -1].values
```

```
#training and testing
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=0)
print(x_train)
```

```
[[7.00e+00 1.50e+02 7.80e+01 ... 3.52e+01 6.92e-01 5.40e+01]
 [4.00e+00 9.70e+01 6.00e+01 ... 2.82e+01 4.43e-01 2.20e+01]
 [0.00e+00 1.65e+02 9.00e+01 ... 5.23e+01 4.27e-01 2.30e+01]
 ...
 [4.00e+00 9.40e+01 6.50e+01 ... 2.47e+01 1.48e-01 2.10e+01]
 [1.10e+01 8.50e+01 7.40e+01 ... 3.01e+01 3.00e-01 3.50e+01]
 [5.00e+00 1.36e+02 8.20e+01 ... 0.00e+00 6.40e-01 6.90e+01]]
```

```
#feature scaling using standard scaler
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train= sc.fit_transform(x_train)
x_test=sc.transform(x_test)
```

```
print(x_train)
```

```
[[ 0.90832902  0.91569367  0.44912368 ... 0.37852648  0.67740401
   1.69955804]
 [ 0.03644676 -0.75182191 -0.47230103 ... -0.50667229 -0.07049698
  -0.96569189]
 [-1.12606292  1.38763205  1.06340683 ... 2.54094063 -0.11855487
  -0.88240283]
 ...
 [ 0.03644676 -0.84620959 -0.21634972 ... -0.94927168 -0.95656442
  -1.04898095]
 [ 2.0708387  -1.12937261  0.24436264 ... -0.26640405 -0.50001442
   0.11706589]
 [ 0.32707418  0.47521786  0.65388473 ... -4.07275877  0.52121586
   2.94889395]]
```

```
from sklearn.tree import DecisionTreeClassifier
model=DecisionTreeClassifier()
model.fit(x_train,y_train)
```

```
DecisionTreeClassifier
```

```
y_pred=model.predict(x_test)
```

```
np.array((y_pred, y_test ))
```

```
array([[1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1,
        1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1,
```

```

0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,
0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1,
1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1,
1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0]])

```

```

from sklearn.metrics import classification_report, confusion_matrix
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.88	0.83	0.86	107
1	0.66	0.74	0.70	47
accuracy			0.81	154
macro avg	0.77	0.79	0.78	154
weighted avg	0.81	0.81	0.81	154

```
print(confusion_matrix(y_test, y_pred))
```

```

[[89 18]
 [12 35]]

```

```
print(data.corr())
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	
	Insulin	BMI	DiabetesPedigreeFunction	\	
Pregnancies	-0.073535	0.017683	-0.033523		
Glucose	0.331357	0.221071	0.137337		
BloodPressure	0.088933	0.281805	0.041265		
SkinThickness	0.436783	0.392573	0.183928		
Insulin	1.000000	0.197859	0.185071		
BMI	0.197859	1.000000	0.140647		
DiabetesPedigreeFunction	0.185071	0.140647	1.000000		
Age	-0.042163	0.036242	0.033561		
Outcome	0.130548	0.292695	0.173844		
	Age	Outcome			
Pregnancies	0.544341	0.221898			
Glucose	0.263514	0.466581			
BloodPressure	0.239528	0.065068			
SkinThickness	-0.113970	0.074752			
Insulin	-0.042163	0.130548			
BMI	0.036242	0.292695			
DiabetesPedigreeFunction	0.033561	0.173844			
Age	1.000000	0.238356			
Outcome	0.238356	1.000000			

```

import seaborn as sns
sns.heatmap(data.corr(), annot=True)

```

&lt;Axes: &gt;



```

from sklearn.ensemble import BaggingClassifier
from sklearn.metrics import accuracy_score
base_model = DecisionTreeClassifier(random_state=42)
bagging_model_dt = BaggingClassifier(base_model, n_estimators=10, random_state=42)
# Train and evaluate the bagging classifier with decision trees
bagging_model_dt.fit(x_train, y_train)
y_pred_dt = bagging_model_dt.predict(x_test)
accuracy_dt = accuracy_score(y_test, y_pred_dt)

```

```
print("Bagging Classifier with Decision Trees Accuracy:", accuracy_dt)
```

```
Bagging Classifier with Decision Trees Accuracy: 0.7857142857142857
```

## 2nd data accuracy

```

x=data.iloc[:, :-3].values
y=data.iloc[:, -1].values

```

```

#training and testing
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=0)
print(x_train)

```

```

[[ 7. 150.  78.  29. 126.  35.2]
 [ 4.  97.  60.  23.   0.  28.2]
 [ 0. 165.  90.  33. 680.  52.3]
 ...
 [ 4.  94.  65.  22.   0.  24.7]
 [11.  85.  74.   0.   0.  30.1]
 [ 5. 136.  82.   0.   0.   0.  ]

```

```

#feature scaling using standard scaler
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train= sc.fit_transform(x_train)
x_test=sc.transform(x_test)

```

```

from sklearn.tree import DecisionTreeClassifier
model=DecisionTreeClassifier()
model.fit(x_train,y_train)

```

```
DecisionTreeClassifier
DecisionTreeClassifier()
```

```
y_pred=model.predict(x_test)
```

```
from sklearn.metrics import classification_report, confusion_matrix
print(classification_report(y_test, y_pred))
```

```

              precision    recall  f1-score   support

     0       0.83        0.82        0.83        107
     1       0.60        0.62        0.61         47

 accuracy          0.76
 macro avg         0.72
 weighted avg      0.76
```


```
print(confusion_matrix(y_test, y_pred))
```

```
[[88 19]
 [18 29]]
```

```
import seaborn as sns
sns.heatmap(data.corr(), annot=True)
```



<Axes: >



```

from sklearn.ensemble import BaggingClassifier
from sklearn.metrics import accuracy_score
base_model = DecisionTreeClassifier(random_state=42)
bagging_model_dt = BaggingClassifier(base_model, n_estimators=10, random_state=42)
#classifier with decision trees
bagging_model_dt.fit(x_train, y_train)
y_pred_dt = bagging_model_dt.predict(x_test)
accuracy_dt = accuracy_score(y_test, y_pred_dt)

print("Bagging Classifier with DT Accuracy:", accuracy_dt)

Bagging Classifier with DT Accuracy: 0.7662337662337663
```